

Transfer learning (ResNet18)

Imports + device

```
import os
import torch
import pandas as pd
from PIL import Image
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms
from sklearn.model_selection import train_test_split

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(device)

cuda
```

Image transforms

Two separate transform pipelines one for the training images, and one for the validation images, mostly the same, but with training images, we augment them, with ColorJitter and RandomRotation to reduce overfitting. At the end we normalize the image to ResNet18 expected input normalization, this is important because pretrained models expect input data scaled to match the distribution they were trained on, and for faster convergence.

```
IMG_SIZE = 64

train_transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ColorJitter(brightness=0.3, contrast=0.3),
    transforms.RandomRotation(10),
    transforms.ToTensor(), # Normalizing each pixel value between 0-1
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])),
])

val_transform = transforms.Compose([
    transforms.Resize((IMG_SIZE, IMG_SIZE)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])),
])
```

Load CSV + train/validate split

Here we load the CSV data and create 85/15 train/validate split. We use stratification to pick the images, so that not one class is "over picked".

```
data_root = '/content/data'

full_dataframe = pd.read_csv(f'{data_root}/Train.csv')

train_dataframe, val_dataframe = train_test_split(
    full_dataframe,
    test_size=0.15,
    stratify=full_dataframe['ClassId'],
    random_state=42
)

print(len(train_dataframe), len(val_dataframe))

33327 5882
```

The Dataset class

Pytorch DataLoader needs a class with **len** and **getitem** to know how many samples exist and how to fetch one by index, this is what our class provides. It takes a dataframe, the image root folder, and a transform as parameters, so the same class can be used for training, validation and test sets, just by passing in different dataframes/transforms.

```
class GTSRBDataset(Dataset):
    def __init__(self, dataframe, root, transform):
        self.dataframe = dataframe.reset_index(drop=True)
        self.root = root
        self.transform = transform

    def __len__(self):
        return len(self.dataframe)

    def __getitem__(self, idx):
        row = self.dataframe.iloc[idx]
        image = Image.open(os.path.join(self.root,
row['Path'])).convert('RGB')
        image = self.transform(image)
        label = int(row['ClassId'])
        return image, label
```

Creating the Dataset objects and DataLoaders

Here we instantiate the class we made before, one for the train set and one for the validation set each with its own dataframe and transform. Each Dataset is then wrapped in a DataLoader

which handles batching, and parallel image loading. Training data is shuffled each epoch so the model doesn't pick up on any accidental ordering in the data.

```
BATCH_SIZE = 64

train_dataset = GTSRBDataset(train_dataframe, data_root,
                              train_transform)
val_dataset = GTSRBDataset(val_dataframe, data_root, val_transform)

train_loader = DataLoader(train_dataset, batch_size=BATCH_SIZE,
                           shuffle=True, num_workers=2)
val_loader = DataLoader(val_dataset, batch_size=BATCH_SIZE,
                        shuffle=False, num_workers=2)

print(len(train_dataset), len(val_dataset))

33327 5882
```

Loading pretrained ResNet18 + swapping the final layer

Here we load the pretrained resnet18, and we swap the final fully connected layer with our own, linear layer, input size is the same as the original last layer input size, we just change the output to our number of classes.

ResNet18 comes pretrained on 1.2M images with 1000 classes, so its convolutional layers already know how to detect general visual features (edges, shapes, textures), these are features that transfer well to new image tasks.

```
import torch.nn as nn
from torchvision import models

NUM_CLASSES = 43

model = models.resnet18(weights='IMAGENET1K_V1')
model.fc = nn.Linear(model.fc.in_features, NUM_CLASSES) # Replacing
the last layer with a brand new linear layer, same number inputs, 43
output instead of 1000
model = model.to(device)
```

Loss function + optimizer

CrossEntropyLoss measures how far the predicted class scores are from the correct label. Adam updates the model's weights each step to reduce that loss, using a small learning rate, since we're fine-tuning pretrained weights and want to nudge them just slightly, rather than overwrite what they have already learned.

```
import torch.optim as optim
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-4)
```

Training loop setup

Creates a weights/resnet18 folder to store checkpoints, sets the number of training epochs, and initialized best_val_accuracy to 0, so that we can track the best performing model as the training progresses.

```
import os

os.makedirs('weights/resnet18', exist_ok=True)

EPOCHS = 15
best_val_accuracy = 0.0
```

Training step (one epoch)

Runs one full pass over the training data. model.train() puts the model into training mode. for each batch: gradients are cleared (optimizer.zero_grad()), a forward pass produces predictions (model(images)) the loss is computed againsts the true labels, loss.backward() computes gradients with backpropagation, and optimizer.step() updates the modal's weights using these gradients. Running loss and accuracy are averaged across all batches (at the end of an epoch).

```
def train_one_epoch():
    model.train()
    running_loss = 0.0
    correct = 0
    total = 0

    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        outputs = model(images)
        loss = criterion(outputs, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        _, predicted = torch.max(outputs, 1)
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

    epoch_loss = running_loss / total
    epoch_accuracy = correct / total
    return epoch_loss, epoch_accuracy
```

Validation step

Runs one full pass over the validation data to check how well the model generalizes to data it wasn't trained on. `model.eval()` switches the model to evaluation mode (dropout/batchnorm behave differently than in training - dropout turns off, batchnorm uses running statistics instead of the current batch). `torch.no_grad()` disabled gradient tracking entirely, since we're only measuring performance here, not updating the weights, this also makes it faster and uses less memory. Otherwise it mirrors `train_one_epoch()`, just no `zero_grad()`, `backward()` or `optimizer.step()`, since nothing is being learned.

```
def validate_one_epoch():
    model.eval()
    running_loss = 0.0
    correct = 0
    total = 0

    with torch.no_grad():
        for images, labels in val_loader:
            images, labels = images.to(device), labels.to(device)

            outputs = model(images)
            loss = criterion(outputs, labels)

            running_loss += loss.item() * images.size(0)
            _, predicted = torch.max(outputs, 1)
            correct += (predicted == labels).sum().item()
            total += labels.size(0)

    epoch_loss = running_loss / total
    epoch_accuracy = correct / total
    return epoch_loss, epoch_accuracy
```

Main training loop

Runs the full training loop for 15 epochs. Each epoch: `train_one_epoch()` trains on the training set, `validate_one_epoch()` checks performance on the validation step, and the results (loss/accuracy for both, plus how long each epoch took) are printed to track progress over time.

Two checkpoints are saved each epoch: `last.pt` is overwritten every single epoch, always holding the most recent model state. `best.pt` is only overwritten when validation accuracy improves on the previous best.

```
import time

for epoch in range(EPOCHS):
    start_time = time.time()

    train_loss, train_accuracy = train_one_epoch()
    val_loss, val_accuracy = validate_one_epoch()
```

```

epoch_time = time.time() - start_time

print(f'Epoch {epoch+1}/{EPOCHS} | '
      f'Train Loss: {train_loss:.4f} Acc: {train_accuracy:.4f} | '
      f'Val Loss: {val_loss:.4f} Acc: {val_accuracy:.4f} | '
      f'Time: {epoch_time:.1f}s')

torch.save(model.state_dict(), 'weights/resnet18/last.pt')

if val_accuracy > best_val_accuracy:
    best_val_accuracy = val_accuracy
    torch.save(model.state_dict(), 'weights/resnet18/best.pt')
    print(f' -> New best model saved (val acc:
{val_accuracy:.4f})')

```

Epoch 1/15 | Train Loss: 0.0127 Acc: 0.9965 | Val Loss: 0.0088 Acc: 0.9971 | Time: 45.0s

Epoch 2/15 | Train Loss: 0.0134 Acc: 0.9961 | Val Loss: 0.0047 Acc: 0.9988 | Time: 44.6s

-> New best model saved (val acc: 0.9988)

Epoch 3/15 | Train Loss: 0.0118 Acc: 0.9968 | Val Loss: 0.0071 Acc: 0.9976 | Time: 45.5s

Epoch 4/15 | Train Loss: 0.0094 Acc: 0.9974 | Val Loss: 0.0084 Acc: 0.9973 | Time: 44.6s

Epoch 5/15 | Train Loss: 0.0133 Acc: 0.9960 | Val Loss: 0.0074 Acc: 0.9983 | Time: 44.5s

Epoch 6/15 | Train Loss: 0.0091 Acc: 0.9977 | Val Loss: 0.0077 Acc: 0.9973 | Time: 47.6s

Epoch 7/15 | Train Loss: 0.0087 Acc: 0.9976 | Val Loss: 0.0141 Acc: 0.9964 | Time: 45.9s

Epoch 8/15 | Train Loss: 0.0077 Acc: 0.9980 | Val Loss: 0.0021 Acc: 0.9995 | Time: 44.5s

-> New best model saved (val acc: 0.9995)

Epoch 9/15 | Train Loss: 0.0074 Acc: 0.9978 | Val Loss: 0.0077 Acc: 0.9978 | Time: 44.2s

Epoch 10/15 | Train Loss: 0.0093 Acc: 0.9974 | Val Loss: 0.0047 Acc: 0.9985 | Time: 44.4s

Epoch 11/15 | Train Loss: 0.0076 Acc: 0.9978 | Val Loss: 0.0022 Acc: 0.9995 | Time: 44.5s

Epoch 12/15 | Train Loss: 0.0038 Acc: 0.9989 | Val Loss: 0.0022 Acc: 0.9993 | Time: 46.9s

Epoch 13/15 | Train Loss: 0.0051 Acc: 0.9987 | Val Loss: 0.0144 Acc: 0.9968 | Time: 46.3s

Epoch 14/15 | Train Loss: 0.0078 Acc: 0.9979 | Val Loss: 0.0027 Acc: 0.9995 | Time: 46.4s

Epoch 15/15 | Train Loss: 0.0073 Acc: 0.9982 | Val Loss: 0.0027 Acc: 0.9991 | Time: 44.2s

Custom CNN

Custom CNN Class skeleton + conv layers

A CNN without pretrained weights compare againsts ResNet18 transfer-learning baseline. It has 3 convolutional blocks, each doubling the channel count (3 -> 32 -> 64 -> 128) while MaxPool2d halves the spatial size after each block (64 -> 32 -> 16 -> 8). Each block is Conv2d -> BatchNorm2d -> ReLU -> MaxPool: the conv layer extract features, batchnorm stabilizes/normalizes activations for more stable training, and ReLU adds non-linearity.

After the 3rd block, the output (128 channels of 8x8 feature maps) is flattened and passed through two fully connected layers (fc1: 8192 -> 256, fc2: 256 -> 43) to produce final class scores. Dropout(0.5) between them randomly zeroes out half the activations during training only, to reduce overfitting.

```
class CustomCNN(nn.Module):
    def __init__(self, num_classes):
        super().__init__()

        self.conv1 = nn.Conv2d(3, 32, kernel_size=3, padding=1)
        self.bn1 = nn.BatchNorm2d(32)

        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, padding=1)
        self.bn2 = nn.BatchNorm2d(64)

        self.conv3 = nn.Conv2d(64, 128, kernel_size=3, padding=1)
        self.bn3 = nn.BatchNorm2d(128)

        self.pool = nn.MaxPool2d(2, 2)
        self.relu = nn.ReLU()

        self.fc1 = nn.Linear(128 * 8 * 8, 256)
        self.dropout = nn.Dropout(0.5)
        self.fc2 = nn.Linear(256, num_classes)

    def forward(self, x):
        x = self.pool(self.relu(self.bn1(self.conv1(x))))
        x = self.pool(self.relu(self.bn2(self.conv2(x))))
        x = self.pool(self.relu(self.bn3(self.conv3(x))))

        x = x.view(x.size(0), -1) # Converting shape to 1d for each
        # batch, from (batch_size, 128, 8, 8) to (batch_size, 128 * 8 * 8)

        x = self.relu(self.fc1(x))
        x = self.dropout(x) # Randomly zeoring 50% of the 256
        # activations
        x = self.fc2(x) # Final layer

        return x
```

Instantiate the model + optimizer

Creating a new CustomCNN instance on GPU. Uses the same CrossEntropyLoss and optimizer, but with a higher learning rate. Since this model starts from scratch it needs larger update steps to learn useful features.

```
custom_model = CustomCNN(num_classes=43).to(device)

criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(custom_model.parameters(), lr=1e-3)

best_val_accuracy_custom = 0.0

os.makedirs('weights/custom_cnn', exist_ok=True)
```

Training loop

Same structure as the ResNet training loop, reusing the exact same train_one_epoch()/validate_one_epoch() functions rather than rewriting them.

```
model = custom_model  # so train_one_epoch/validate_one_epoch operate
                        on the custom CNN

EPOCHS = 15

for epoch in range(EPOCHS):
    start_time = time.time()

    train_loss, train_accuracy = train_one_epoch()
    val_loss, val_accuracy = validate_one_epoch()

    epoch_time = time.time() - start_time

    print(f'Epoch {epoch+1}/{EPOCHS} | '
          f'Train Loss: {train_loss:.4f} Acc: {train_accuracy:.4f} | '
          f'Val Loss: {val_loss:.4f} Acc: {val_accuracy:.4f} | '
          f'Time: {epoch_time:.1f}s')

    torch.save(model.state_dict(), 'weights/custom_cnn/last.pt')

    if val_accuracy > best_val_accuracy_custom:
        best_val_accuracy_custom = val_accuracy
        torch.save(model.state_dict(), 'weights/custom_cnn/best.pt')
        print(f' -> New best model saved (val acc:
{val_accuracy:.4f})')

Epoch 1/15 | Train Loss: 2.1268 Acc: 0.3605 | Val Loss: 0.8699 Acc:
0.7112 | Time: 43.2s
-> New best model saved (val acc: 0.7112)
```

```
Epoch 2/15 | Train Loss: 1.0008 Acc: 0.6575 | Val Loss: 0.2919 Acc: 0.9104 | Time: 40.3s
-> New best model saved (val acc: 0.9104)
Epoch 3/15 | Train Loss: 0.6883 Acc: 0.7606 | Val Loss: 0.1159 Acc: 0.9728 | Time: 41.6s
-> New best model saved (val acc: 0.9728)
Epoch 4/15 | Train Loss: 0.5490 Acc: 0.8081 | Val Loss: 0.0698 Acc: 0.9847 | Time: 41.3s
-> New best model saved (val acc: 0.9847)
Epoch 5/15 | Train Loss: 0.4796 Acc: 0.8292 | Val Loss: 0.0638 Acc: 0.9862 | Time: 41.8s
-> New best model saved (val acc: 0.9862)
Epoch 6/15 | Train Loss: 0.4358 Acc: 0.8438 | Val Loss: 0.0780 Acc: 0.9857 | Time: 40.3s
Epoch 7/15 | Train Loss: 0.4024 Acc: 0.8550 | Val Loss: 0.0464 Acc: 0.9903 | Time: 41.6s
-> New best model saved (val acc: 0.9903)
Epoch 8/15 | Train Loss: 0.3645 Acc: 0.8675 | Val Loss: 0.0344 Acc: 0.9923 | Time: 41.6s
-> New best model saved (val acc: 0.9923)
Epoch 9/15 | Train Loss: 0.3442 Acc: 0.8771 | Val Loss: 0.0304 Acc: 0.9944 | Time: 41.7s
-> New best model saved (val acc: 0.9944)
Epoch 10/15 | Train Loss: 0.3284 Acc: 0.8795 | Val Loss: 0.0223 Acc: 0.9949 | Time: 39.4s
-> New best model saved (val acc: 0.9949)
Epoch 11/15 | Train Loss: 0.3151 Acc: 0.8845 | Val Loss: 0.0245 Acc: 0.9952 | Time: 41.8s
-> New best model saved (val acc: 0.9952)
Epoch 12/15 | Train Loss: 0.3047 Acc: 0.8898 | Val Loss: 0.0227 Acc: 0.9944 | Time: 41.2s
Epoch 13/15 | Train Loss: 0.2912 Acc: 0.8921 | Val Loss: 0.0169 Acc: 0.9963 | Time: 41.5s
-> New best model saved (val acc: 0.9963)
Epoch 14/15 | Train Loss: 0.2800 Acc: 0.8972 | Val Loss: 0.0191 Acc: 0.9954 | Time: 39.5s
Epoch 15/15 | Train Loss: 0.2693 Acc: 0.9004 | Val Loss: 0.0229 Acc: 0.9952 | Time: 41.7s
```

Evaluation and analysis

Loading Test.csv + building test Dataset/DataLoader

```
test_dataframe = pd.read_csv(f'{data_root}/Test.csv')
print(test_dataframe.shape)

test_dataset = GTSRBDataset(test_dataframe, data_root, val_transform)
```

```
test_loader = DataLoader(test_dataset, batch_size=BATCH_SIZE,
shuffle=False, num_workers=2)
```

```
print(len(test_dataset))
```

```
(12630, 8)
12630
```

Loading both models from their best.pt checkpoints

```
resnet_model = models.resnet18(weights=None)
resnet_model.fc = nn.Linear(resnet_model.fc.in_features, NUM_CLASSES)
resnet_model.load_state_dict(torch.load('weights/resnet18/best.pt'))
resnet_model = resnet_model.to(device)
resnet_model.eval()
```

```
custom_model_eval = CustomCNN(num_classes=43)
custom_model_eval.load_state_dict(torch.load('weights/custom_cnn/best.pt'))
custom_model_eval = custom_model_eval.to(device)
custom_model_eval.eval()
```

```
CustomCNN(
  (conv1): Conv2d(3, 32, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (bn1): BatchNorm2d(32, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv2): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (conv3): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1),
padding=(1, 1))
  (bn3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
  (relu): ReLU()
  (fc1): Linear(in_features=8192, out_features=256, bias=True)
  (dropout): Dropout(p=0.5, inplace=False)
  (fc2): Linear(in_features=256, out_features=43, bias=True)
)
```

Evaluation loop

```
def evaluate(model, loader):
    model.eval()
    correct = 0
    total = 0
```

```

all_preds = []
all_labels = []

with torch.no_grad():
    for images, labels in loader:
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = torch.max(outputs, 1)

        correct += (predicted == labels).sum().item()
        total += labels.size(0)

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

accuracy = correct / total
return accuracy, all_preds, all_labels

resnet_test_acc, resnet_preds, resnet_labels = evaluate(resnet_model,
test_loader)
custom_test_acc, custom_preds, custom_labels =
evaluate(custom_model_eval, test_loader)

print(f'ResNet18 Test Accuracy: {resnet_test_acc:.4f}')
print(f'Custom CNN Test Accuracy: {custom_test_acc:.4f}')

ResNet18 Test Accuracy: 0.9810
Custom CNN Test Accuracy: 0.9766

```

Parameter count comparison

```

resnet_params = sum(p.numel() for p in resnet_model.parameters())
custom_params = sum(p.numel() for p in custom_model_eval.parameters())

print(f'ResNet18 parameters: {resnet_params:,}')
print(f'Custom CNN parameters: {custom_params:,}')
print(f'ResNet18 is {resnet_params / custom_params:.1f}x larger')

ResNet18 parameters: 11,198,571
Custom CNN parameters: 2,202,155
ResNet18 is 5.1x larger

```

Confusion matrix (top confusion)

```

from sklearn.metrics import confusion_matrix

resnet_cm = confusion_matrix(resnet_labels, resnet_preds)
custom_cm = confusion_matrix(custom_labels, custom_preds)

print(resnet_cm.shape)

```

(43, 43)

```
import numpy as np

def top_confusions(cm, n=10):
    cm_errors = cm.copy()
    np.fill_diagonal(cm_errors, 0)  # zero out correct predictions,
    keep only mistakes

    confusions = []
    for true_class in range(cm_errors.shape[0]):
        for pred_class in range(cm_errors.shape[1]):
            count = cm_errors[true_class, pred_class]
            if count > 0:
                confusions.append((true_class, pred_class, count))

    confusions.sort(key=lambda x: x[2], reverse=True)
    return confusions[:n]

print('ResNet18 top confusions (true -> predicted, count):')
print(top_confusions(resnet_cm))

print('\nCustom CNN top confusions (true -> predicted, count):')
print(top_confusions(custom_cm))

ResNet18 top confusions (true -> predicted, count):
[(25, 27, np.int64(25)), (4, 2, np.int64(12)), (27, 28, np.int64(12)),
 (38, 40, np.int64(11)), (18, 26, np.int64(10)), (21, 11, np.int64(9)),
 (15, 3, np.int64(8)), (27, 24, np.int64(8)), (9, 3, np.int64(7)), (22,
 25, np.int64(7))]

Custom CNN top confusions (true -> predicted, count):
[(27, 24, np.int64(30)), (12, 32, np.int64(28)), (8, 5, np.int64(25)),
 (3, 5, np.int64(21)), (21, 11, np.int64(21)), (1, 2, np.int64(15)),
 (26, 22, np.int64(14)), (22, 25, np.int64(10)), (25, 20, np.int64(9)),
 (31, 5, np.int64(6))]
```

Misclassified Examples

```
class_names = {
    0: 'Speed limit (20km/h)', 1: 'Speed limit (30km/h)', 2: 'Speed
    limit (50km/h)',
    3: 'Speed limit (60km/h)', 4: 'Speed limit (70km/h)', 5: 'Speed
    limit (80km/h)',
    6: 'End of speed limit (80km/h)', 7: 'Speed limit (100km/h)', 8:
    'Speed limit (120km/h)',
    9: 'No passing', 10: 'No passing for vehicles over 3.5 tons',
    11: 'Right-of-way at next intersection', 12: 'Priority road', 13:
    'Yield',
    14: 'Stop', 15: 'No vehicles', 16: 'Vehicles over 3.5 tons'
```

```

prohibited',
    17: 'No entry', 18: 'General caution', 19: 'Dangerous curve left',
    20: 'Dangerous curve right', 21: 'Double curve', 22: 'Bumpy road',
    23: 'Slippery road', 24: 'Road narrows on the right', 25: 'Road
work',
    26: 'Traffic signals', 27: 'Pedestrians', 28: 'Children crossing',
    29: 'Bicycles crossing', 30: 'Beware of ice/snow', 31: 'Wild
animals crossing',
    32: 'End of all speed/passing limits', 33: 'Turn right ahead', 34:
'Turn left ahead',
    35: 'Ahead only', 36: 'Go straight or right', 37: 'Go straight or
left',
    38: 'Keep right', 39: 'Keep left', 40: 'Roundabout mandatory',
    41: 'End of no passing', 42: 'End of no passing (vehicles > 3.5
tons)'
}

```

```

np.random.seed(42)

```

```

def show_misclassified(preds, labels, df, root, model_name, n=6):
    preds = np.array(preds)
    labels = np.array(labels)
    wrong_idx = np.where(preds != labels)[0]
    chosen = np.random.choice(wrong_idx, size=n, replace=False)

    fig, axes = plt.subplots(1, n, figsize=(3 * n, 3))
    fig.suptitle(f'{model_name} - Misclassified Examples')

    for ax, idx in zip(axes, chosen):
        row = df.iloc[idx]
        img = Image.open(os.path.join(root,
row['Path'])).convert('RGB')
        true_name = class_names[labels[idx]]
        pred_name = class_names[preds[idx]]

        ax.imshow(img)
        ax.set_title(f'True: {true_name}\nPred: {pred_name}',
fontsize=8)
        ax.axis('off')

    plt.tight_layout()
    plt.show()

```

```

show_misclassified(resnet_preds, resnet_labels, test_dataframe,
data_root, 'ResNet18', n=6)
show_misclassified(custom_preds, custom_labels, test_dataframe,
data_root, 'Custom CNN', n=6)

```



Convergence graph

```
resnet_val_acc = [0.9971, 0.9988, 0.9976, 0.9973, 0.9983, 0.9973,
0.9964,
0.9995, 0.9978, 0.9985, 0.9995, 0.9993, 0.9968,
0.9995, 0.9991]
custom_val_acc = [0.7112, 0.9104, 0.9728, 0.9847, 0.9862, 0.9857,
0.9903,
0.9923, 0.9944, 0.9949, 0.9952, 0.9944, 0.9963,
0.9954, 0.9952]

epochs = range(1, 16)

plt.figure(figsize=(9, 5))
plt.plot(epochs, resnet_val_acc, marker='o', color='#0072B2',
label='ResNet18 (transfer learning)')
plt.plot(epochs, custom_val_acc, marker='o', color='#E69F00',
label='Custom CNN (from scratch)')

plt.xlabel('Epoch')
plt.ylabel('Validation Accuracy')
plt.title('Validation Accuracy Convergence: ResNet18 vs Custom CNN')
plt.ylim(0, 1.02)
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```

